



Patch-Based Diffusion Models for Solving Inverse Problems in Computed Tomography Imaging

Executive Summary

Thomas Hartigan | MPhil in Data Intensive Science

Purpose and Motivation

Computed tomography (CT) reconstructs anatomy from X-ray measurements. Fewer projection angles can reduce acquisition time and radiation exposure, but leave reconstruction underdetermined because many images explain sparse measurements. Classical methods therefore produce streaking or blur without a prior for plausible anatomy.

Diffusion models provide expressive learned priors, although whole-image models can require extensive data, training and GPU memory. Hu et al. proposed the Patch Diffusion Inverse Solver (PaDIS) as an alternative. PaDIS trains on position-labelled patches and repeatedly changes their partition during reconstruction, combining local scores into a whole-image update intended to remove seams. The original paper reported that PaDIS consistently outperformed a comparable whole-image diffusion model on sparse-view CT.

This project tested whether that result reproduces on a public dataset with a specified fan-beam geometry, and which choices drive performance. Independent testing matters because hidden implementation choices make results difficult to validate or build upon, making this a reproduction and extension rather than a numerical rerun.

Approach

We reimplemented training, tuning and evaluation within LION rather than treating the released code as executable ground truth. Experiments used LIDC-IDRI slices with patient-separated training, validation and test sets. The default subset contained at most four slices per patient, while an ablation used every processed slice. A fully defined two-dimensional fan-beam operator simulated noise-free measurements, isolating angular undersampling from photon-count effects.

Because the paper and code differ in algorithmic choices, we used three tracks to capture different interpretations of the source material. The paper-style track followed the published equations and sampling choices, while the public-compatible track followed the released repository and reproduced an adapted fork's outputs within sampler variation on matched inputs. The LION-physics extension used the same priors but normalised data-consistency gradients using the squared norm of the composed CT operator, avoiding geometry-specific calibration constants.

The evaluation compared classical, whole-image diffusion and PaDIS reconstruction. Main tests used 8 and 20 full-angle views, with extensions to 60 views, limited angles and native 512×512 resolution. Ablations varied patch size, training-set size, positional encoding, initialisation and noise schedule. Results used PSNR, SSIM, MAE, measurement residuals and image-level bootstrap intervals.

Principal Findings

Figure 1 shows representative fan-beam reconstructions from PaDIS and the comparison methods. On the exploratory four-image native 512×512 subset (Figure 2), LION-physics PaDIS achieved 40.08 dB PSNR and 0.960 SSIM, exceeding FDK and CP-TV without requiring a whole-image model at that resolution. The separately tuned PaDIS VE-DDNM configuration produced the best 20-view and 60-view PSNR




Figure 1: 60-view HU (top) and 20-view NI (bottom) reconstructions comparing FDK, CP, whole-image VE-DPS and PaDIS-DPS.

values in the study, reaching 36.49 and 41.55 dB respectively, while the operator-normalised VE-DPS rows reduced relative sinogram residuals to approximately 10^{-3} and therefore kept the outputs closely tied to the measurements.




Figure 2: Native-resolution 60-view reconstruction, displayed in HU.

The stronger claim of consistent PaDIS superiority was not reproduced. Under the VE-DPS comparison, the whole-image model achieved 34.25 dB at 20 views and 29.58 dB at 8 views, compared with 32.82 and 27.65 dB for PaDIS, and the paper-style track showed the same reversal. PaDIS exceeded the whole-image result only for some separately selected sampler configurations, notably VE-DDNM, so these comparisons do not isolate the prior alone. Patch training is therefore viable and sometimes advantageous, but the result depends on the interaction between prior, sampler and forward-operator scaling.

The ablations changed the interpretation further. A 96×96 patch model had almost the same mean PSNR as the whole-image model, while intermediate smaller patches performed worse, suggesting that global context remained valuable. However, the non-default patch and no-position models inherited the default PaDIS inference setting and were not individually optimised. Positional channels improved measurement-conditioned PSNR by only 0.3 to 0.4 dB, in contrast to the original paper’s reported failure without them, and FDK and Gaussian initialisation converged to similar final quality. The geometric noise schedule instead improved PSNR by about 1.3 dB over the released code’s EDM-style default. Using all

available training slices did not improve PaDIS under a fixed training-time budget, although fewer effective epochs mean this is not evidence that additional data are intrinsically unhelpful.

Among the four untuned samples per method in Figure 3, random PaDIS partitions avoided the rectangular boundaries created by fixed patch stitching, although this small set cannot rank distributional quality. The pipeline recorded roughly two minutes per PaDIS slice, compared with about 2.6 minutes for whole-image VE-DPS and 24 minutes for fixed averaging or stitching, but concurrent jobs prevent an isolated speed comparison.

Reproducibility, Limitations and Implications

Exact reproduction from the original paper was not possible. Important architecture information was distributed inconsistently between references, appendix and code. Training patch counts, validation and checkpoint selection, test slice identities, complete physical geometry, and the hyperparameter search protocol were not specified. The paper and released code form a combined but internally inconsistent specification, containing materially different algorithmic choices. The three-track design made these differences explicit and showed that they can change scientific conclusions. The reversal of the central comparison is therefore an informative reproducibility result rather than a failed outcome.

The results remain limited to one dataset, simulated two-dimensional measurements and a small test set. Noise-free sinograms omit photon noise, scatter, beam hardening and motion, so the work does not establish clinical low-dose performance. Some expensive experiments used only four images, confidence intervals do not include retraining variability, and no radiologist assessed lesion preservation or hallucination. A larger native-resolution test set is therefore needed before drawing a population-level comparison between reconstruction methods.

The practical conclusion is that PaDIS should be viewed as a memory-scalable route to high-resolution priors, not as a universally superior replacement for whole-image diffusion. Future work should use equal-compute and equal-epoch comparisons, multiple training seeds, calibrated measurement noise, external datasets and clinically relevant task assessment for safe clinical deployment. The project also demonstrates that inverse-problem papers should publish complete geometry, data splits, tuning procedures and code that matches the stated algorithms. Otherwise, apparent method improvements may instead reflect hidden implementation choices.




Figure 3: Unconditional samples from whole-image, fixed patch-stitching, fixed patch-averaging and PaDIS priors, shown in NI.